

R3 v4

AI-Native Digital Audio Workstation

Product Requirements Document — Version 4.1

Version	4.1 — Engineering-Verified Investor-Grade
Status	Active — MVP Build Phase (Items 1–3 Complete)
Updated	2026-04-12 (v4.1 session update)
Classification	Confidential
Audience	Engineering, Investors, Design, QA
Owner	Ty (Founder / Lead Engineer)
Codebase	~/Stable — pnpm monorepo — feature//lpte-extraction
Deployment	Railway (backend) + Vercel (frontend)
Hygiene Baseline	TSC: 0 errors — 11 routers wired — aiDecisionLog migration 0005

CHANGELOG v4.1 (2026-04-12): MultitrackView double transport bar fixed (/mixer page). Acid-techno theme applied to mixer. SKILLS.md created (22 engineering patterns). DEMO_CHECKLIST.md created. Migration 0005 applied to local DB — Railway production apply in progress. | CHANGELOG FROM PRD v3: LemonSqueezy removed. Tier names corrected to explorer/creator/pro_artist. State management corrected from Redux to Zustand. appRouter shape updated to reflect 11 wired routers. aiDecisionLog table added. SessionChip + SessionSummaryPanel wired. Future Development section added.

0. Executive Summary

R3 v4 is an AI-native Digital Audio Workstation built for the market existing tools have failed: DJs and music creators who need professional-quality results without professional-level technical skill or time investment.

The global music production software market is valued at \$1.4B (2024) and growing at 8.2% CAGR. No dominant player has shipped a genuinely AI-native workflow layer. Ableton, Logic, and Rekordbox remain fundamentally unchanged in their core interaction model for over a decade.

R3's differentiation is architectural, not cosmetic. The Low-Latency Processing Transition Engine (LLPTE) is a named, inspectable, real-time AI pipeline with confirmed sub-15ms inference latency. Every AI decision is transparent, overridable, and quantifiably faster than its manual equivalent.

Current Build State (Verified 2026-04-09)

MVP Item	Status	Evidence
AI Auto-Leveling	✓ Complete	6 layers, 20 Vitest tests, LLPTE wired
Smart Transitions	✓ Complete	9 files, 22 Vitest tests, Camelot scoring
Time Savings Tracking	✓ Wired	SessionChip + SessionSummaryPanel in DAW.tsx
Mix Suggestion System	▣ In Progress	MixSuggestionsPanel.tsx built, backend pending
aiDecisionLog table	✓ Migration 0005	11 columns, Drizzle schema + types
appRouter	✓ 11 routers	sessions, sessionMetrics, admin, daw, subscription, mixer, dj, aiMix, projects, presets, settings
TSC	✓ Zero errors	Verified post all patches
Railway deployment	✓ Live	railway.toml confirmed

One-sentence pitch: R3 makes professional-quality mixing accessible to anyone with a track and a deadline — and proves it with a number on screen.

1. Product Summary & Vision

Identity

Product Name	R3 v4
Category	AI-Native Digital Audio Workstation (DAW)
Build Target	TypeScript pnpm Monorepo
Frontend	React 18 + Vite 5 (client/)
Backend	Express + tRPC + WebSocket (server/)
Database	PostgreSQL via Drizzle ORM 0.39.3
AI Pipeline	LLPTE — 6 monorepo packages (packages/llpte-*)
Rendering	WebGPU + Three.js r128 / React Three Fiber
Auth	bcrypt v6.0.0 + JWT (jsonwebtoken v9.0.3)
Payments	Stripe v20.4.1 ONLY — no LemonSqueezy
State Management	Zustand (NOT Redux — corrected from PRD v3)
Routing	Wouter (NOT react-router-dom)
Testing	Vitest — 42+ cases across LLPTE layers
Deployment	Railway (backend) + Vercel (frontend)
Design System	Tailwind CSS + Radix UI
Tiers	explorer · creator · pro_artist (Stripe only)

Vision Statement

The future of music creation is not a smarter manual interface. It is a system that understands audio as well as the engineer sitting next to you — one that can hear what you cannot, adjust what you would have missed, and do it before you knew it was a problem.

Design Philosophy

- Dark by default. Every background is zinc-950 or darker.
- Thin borders, maximum information. 1px borders at 6% opacity. Zero decorative chrome.
- Neon as signal, not decoration. Cyan = active. Violet = AI. Amber = warning. Red = danger.
- JetBrains Mono for data, Inter for language. BPM, latency, dB — monospace. Labels — Inter.
- Every AI action leaves a visible trace. Ghost knobs, violet overlays, confidence scores.

2. Market Opportunity

Total Addressable Market

Segment	Size	Source
Global Music Production Software	\$1.4B (2024)	Grand View Research
CAGR through 2030	8.2%	Grand View Research
Projected Market (2030)	\$2.3B	Calculated
DJ Software Segment	\$380M	Statista
Content Creator Audio Tools	\$620M	IBISWorld

Serviceable Addressable Market

- Digital DJs: 4.2M active globally (DJPA 2023). 1.8M use software-based tools. Average spend: \$180/year. SAM: \$324M.
- Content Creators: 50M+ active (YouTube, TikTok, Podcasts). 12M spend on audio tools. Average spend: \$85/year. SAM: \$1.02B.
- Combined SAM: approximately \$1.34B.

Why Now

- WebGPU is production-ready. Real-time GPU-accelerated audio visualization in a browser-native desktop application — no plugin required.
- Quantized inference is fast enough. Sub-15ms inference on consumer hardware was not achievable 24 months ago. R3 ships at 10ms.
- The creator economy demands professional output. Audiences trained by Spotify and Apple Music reject amateur mixes. R3 solves this.

3. Competitive Landscape

Direct Competitors

Dimension	Ableton Live 12	Traktor Pro 4	R3 v4
AI Integration	None (Max for Live only)	None	Native LLPTE pipeline
Learning Curve	6–12 months	3–6 months	Target: 30 minutes
Price	\$99–\$749 one-time	\$99 one-time	\$20/month (SaaS)
AI Transparency	N/A	N/A	Ghost knobs + confidence scores
Time Savings	None	None	Native — session-level tracking
Platform	Desktop only	Desktop only	Desktop-first, web-ready

Competitive Moat

R3's moat is not the AI model. The moat is the LLPTE architecture: a named, versioned, inspectable real-time pipeline embedded in the product's visual identity. Competitors cannot copy the node graph, the ghost knobs, or the Time Savings panel without rebuilding their entire product architecture.

Secondary moat — data flywheel: Every accepted AI suggestion is a labeled training signal. At 5,000 users with 3 sessions/week, R3 accumulates approximately 780,000 labeled audio engineering decisions per month. This dataset has no competitor equivalent.

4. Business Model & Monetization

Subscription Tiers (Stripe Only — Corrected from PRD v3)

Note: PRD v3 incorrectly listed tier names as Starter/Pro/Studio. Correct names in codebase and Stripe: explorer / creator / pro_artist.

Tier (DB value)	Price	Target User
explorer	Free	First-time creators, conversion target
creator	\$20/month	Primary commercial tier — DJs and serious creators
pro_artist	\$60/month (Post-MVP)	Professional producers and small studios
Enterprise / API	Custom Pricing (Post-MVP)	Platforms, labels, hardware manufacturers

Revenue Projections (Conservative)

Quarter	Paying Users	MRR	ARR
Q1 Post-Launch	500	\$10,000	\$120,000
Q2	1,500	\$30,000	\$360,000
Q3	3,500	\$70,000	\$840,000
Q4	6,000	\$120,000	\$1,440,000
Year 2 Q4	25,000	\$500,000	\$6,000,000

Unit Economics

CAC (content-led)	≤\$35
LTV at \$20/month	\$240 (12-month avg retention)
LTV:CAC Ratio	≥7:1
Gross Margin	82%
Monthly Churn Target	≤4%
Payback Period	≤2 months

5. Target Users — Deep Personas

Persona 1: Marcus — The Working DJ

Age / Location	28, Chicago
Income	\$200–\$500/night, 3–4 nights/week
Tools Today	Rekordbox + Pioneer CDJ + Ableton for prep
Core Pain	Set prep takes 8 hours. Should take 4.
Willingness to Pay	\$20/month without hesitation if transitions are provably better
Acquisition Channel	r/DJs, DJTT, short-form video

Persona 2: Priya — The Content Creator

Age / Location	24, Austin
Audience	180K YouTube subscribers, lifestyle/travel
Audio Skill	Near zero — treats audio as afterthought
Core Pain	Gets comments about audio quality. Cannot fix it.
Willingness to Pay	\$20/month if it saves 1 hour/week (\$100+ hourly value)
Acquisition Channel	Creator economy newsletters, "pro audio in 10 min" content

Persona 3: Deon — The Bedroom Producer

Age / Location	19, Atlanta
Platform	FL Studio, SoundCloud, seeking A&R placements
Audio Skill	High for sound design, low for mixing/mastering
Core Pain	Mixes are "muddy." Knows it. Cannot fix it.
Willingness to Pay	\$20/month if it gets him a placement or serious mix compliment
Acquisition Channel	SoundCloud community, beat-making subreddits, TikTok

6. Objectives & Success Criteria

Strategic Objectives (Ranked)

- Demonstrate that AI produces a measurably better mix than a human in less time — proven by the Time Savings panel and before/after audio comparison.
- Establish LLPTE as a named, defensible, visible AI architecture that differentiates R3 from every competitor.
- Achieve 65+ NPS among alpha users within 60 days of private launch.
- Reach \$120K ARR within 12 months of public launch.
- Collect 1M+ labeled AI decision signals for model improvement within 18 months.

Feature-Level Success Criteria

Feature	Metric	Target	Status
AI Auto-Leveling	Acceptance rate	≥65%	✔ Built — tracking via aiDecisionLog
AI Auto-Leveling	Inference latency	≤15ms	✔ 10ms confirmed
Smart Transitions	Setup time reduction	≥50%	✔ Built
Time Savings	Panel adoption	≥80% of sessions	✔ Wired in DAW.tsx
Mix Suggestion	Acceptance rate	≥40%	❑ MVP item 4 pending
Overall	Day 7 retention	≥60%	❑ Pre-launch
Overall	Day 30 retention	≥40%	❑ Pre-launch

7. UI Architecture — Five Zones (Verified)

Zone 1 — Top Navigation Bar (48px, full-width)

- Background: absolute black. Zero transparency.
- Left: R3 wordmark — JetBrains Mono, weight 700, letter-spacing 0.2em, pure white.
- Center: Transport controls — skip-back, stop, record, play (cyan glow #00F5FF when active), skip-forward.
- Center: BPM counter — JetBrains Mono, 2 decimal places, tap-tempo adjacent.
- Right: CPU/memory meter, user avatar, tier badge (Pro = violet pill), settings gear.
- Right: SessionChip — shows "● Live" in violet during active session, AI action count. (WIRED 2026-04-09)

Zone 2 — Left Sidebar — Project & Sample Browser (240px)

- Background: zinc-900. Right border: 1px zinc-800.
- Search field — debounced 200ms — queries tRPC sampleSearch procedure.
- File tree: PROJECTS, SAMPLES, MIDI, PRESETS, UPLOADS — chevron collapse/expand.
- Mini waveform preview on hover — 48px height, electric blue #3B82F6.

Zone 3 — Center Arrangement / Timeline (~60% width)

- Background: zinc-950. Playhead: 1px solid #00F5FF with triangle handle and 4px radial feather.
- Eight tracks: KICK, 808 BASS, SYNTH LEAD, CHORD PAD, VOCAL CHOP, FX SEND, AI MIX (LLPTE overlay), MASTER.
- Right-click context menu on clips: Rename, Duplicate, Delete, Bounce to Audio, Send to AI.
- "Send to AI" triggers LLPTE pipeline — violet overlay renders — transition selector appears.

Zone 4A — DJ Mixer / Channel Strips (bottom-left, 220px tall)

- Six vertical channel strips with 3-band EQ, gain trim knob, VU meter, fader.
- Ghost knobs appear when AI Auto-Leveling is active — translucent violet overlay.
- DJ crossfader with AI-suggested position indicator during transitions.

Zone 4B — LLPTE Node Graph (bottom-right)

- Five animated nodes: inputRouter → spectralAnalyzer → llpte-core → transitionGraph → outputBus.
- Connector lines: 1px animated dashes, 5px dash / 5px gap, 30px/second signal direction.
- LLPTE Core: 2× size, arc spinner, inference 10ms badge, pulsing ring during inference.
- Transition Graph tooltip: "847 active edges | last tick 0.8ms" — 200ms hover delay.

- Signal oscilloscope: violet (pre-AI) and cyan (post-AI), 30fps, 40px height.

Zone 5 — Right Sidebar — Effects Rack (280px)

- Six modules: Reverb, Delay, Compressor, Transition Graph (LLPTE), Output Bus (LLPTE), EQ.
- Each module: power toggle, drag handle, expand arrow, active indicator (emerald dot).

8. Core Features — Full Specification

8.1 AI Auto-Leveling — Complete

Gain staging is the most common source of amateur-sounding mixes. R3 addresses it automatically within the first 2 bars of playback via the LLPTE pipeline.

Confidence Gating Logic

confidence $\geq$ 0.65 → auto-apply via `AudioParam.setTargetAtTime()`
confidence $\geq$ 0.40 → surface as ghost knob suggestion only
confidence $<$ 0.40 → discard silently, log to `aiDecisionLog`

8.2 Smart Transitions — Complete

Seven transition types precomputed by `transitionGraph` before the phrase boundary arrives.

Type	Mechanism	Best For
Crossfade (linear)	Equal-power gain crossfade	Same-key tracks
Crossfade (S-curve)	Sigmoidal gain curve	Energy preservation
Filter Sweep	LPF close → HPF open on incoming	High-energy drops
Reverb Tail	Outgoing reverb extended, incoming fades in	Atmospheric sections
Beat Drop Alignment	Transition delayed to nearest downbeat	Drop-centric sets
Echo Freeze	Outgoing pitch-locked, incoming crossfades	Dramatic transitions
Key-Match Crossfade	Camelot wheel harmonic scoring	Maximum harmonic quality

8.3 Time Savings Tracking — Wired (SessionChip + SessionSummaryPanel)

Every session tracked against manual baseline. Displayed in top nav chip and full-screen summary overlay on session end.

Action Type	Manual Baseline	AI Execution
Gain adjustment (one track)	45 seconds	0.5 seconds
EQ sweep (one track)	90 seconds	1 second
Crossfade transition setup	180 seconds	5 seconds
Frequency conflict	300 seconds	0.1 seconds

detection		
-----------	--	--

8.4 Mix Suggestion System — In Progress (MVP Item 4)

MixSuggestionsPanel.tsx component exists. Trigger detection logic in server/services/. Backend wiring to sessionMetrics.router pending.

- Frequency masking detection: FFT overlap in shared bands > -12dBFS for ≥1 bar.
- Dynamic imbalance: RMS differential > 6dB between any two active tracks.
- Arrangement conflict: ≥3 tracks simultaneously peaking above -6dBFS.

8.5 LLPTE Core — Real-Time AI Engine

Package Architecture

packages/llpte-core/	– Pipeline orchestration, lifecycle
packages/llpte-adapters/	– inputRouter, format normalization
packages/llpte-signal/	– spectralAnalyzer, FFT/RMS/LUFS
packages/llpte-ai/	– aiMixEngine, inference, confidence
packages/llpte-transition-graph/	– transitionGraph, curve precompute
packages/llpte-execution/	– outputBus, signal delivery

Performance Contract (SLA)

Metric	Hard Limit	Current (Verified)
Inference latency (p50)	≤15ms	10ms
Inference latency (p99)	≤25ms	TBD
Node tick time	≤1ms	0.8ms
Active edges (MVP)	≤2000	847
Pipeline uptime	99.9% per session	Auto-restart on breach

9. User Flows

Primary Flow: DJ Set Preparation

- Entry: User opens R3, lands on arrangement view (route: /instrument → /daw).
- Step 1 — Track load: drag tracks from sidebar, LLPTE pipeline initializes.
- Step 2 — First playback: Space bar, playhead starts, LLPTE activates within 2 bars.
- Step 3 — Auto-leveling: Toast fires "AI Mix updated — confidence 94%", ghost knobs appear.
- Step 4 — Frequency conflict: Suggestion panel fires, user accepts EQ adjustment.
- Step 5 — Transition zone: Violet overlay appears at bar 15, user selects Key-Match Crossfade.
- Step 6 — Session end: Space bar, SessionSummaryPanel appears, user exports PNG.

Routing Contract (Non-Negotiable)

Post-login redirect	/instrument ONLY — never /daw
User journey	Pricing → Login → Instrument → DAW → Loopstation
ProtectedRoute	Must NOT call hydrateFromToken() during render

10. Non-Goals (Firm Scope Boundary)

Non-Goal	Why Excluded	When Revisited
VST/AU/AAX plugin support	Requires native binary bridge — months of work	Post-Series A
Advanced MIDI composition	Separate product surface	Roadmap Phase 3
Real-time collaboration	Requires conflict resolution logic	Roadmap Phase 2
Mobile platform	Requires AudioWorklet + SharedArrayBuffer (desktop APIs)	Post \$5M ARR
Sound design / synthesis	Different user and use case	Never (separate product)
Offline mode	Cloud inference preferred for model updates	Roadmap Phase 2

11. Technical Requirements & Architecture

Performance Targets

Layer	Metric	Requirement
Audio engine	Round-trip latency	≤10ms
LLPTE inference	p50 latency	≤15ms (current: 10ms)
LLPTE inference	p99 latency	≤25ms
Node tick	Per-node	≤1ms (current: 0.8ms)
UI frame rate	During active playback	≥60fps
WebSocket round-trip	State sync	≤50ms
Database query	Session write	≤100ms
Memory ceiling	8 tracks × 5 min	≤2GB RAM

Stack (Pinned Versions — Verified)

TypeScript	5.9.3
drizzle-orm	0.39.3
ws	8.20.0
Three.js	0.128.0 (r128)
Stripe	20.4.1
Express	4.22.1
Zod	3.25.76
Vitest	Latest — coverage via v8
pnpm	10.33.0
Node	22.x (Active LTS — system confirmed v22.22.0)

HARD GUARDS — Non-Negotiable (CLAUDE.md)

- No any — use unknown + type guard. (15 violations fixed 2026-04-09. 5 remaining: routes/presets.ts × 4, audio-analysis.ts × 1 with eslint-disable exemption — no .d.ts for node-web-audio-api)
- No swallowed exceptions — all async functions handle errors explicitly.
- No console.log in committed code. (5 violations in server/index.ts:300-308 — pending fix)
- No write without read first (Wire.txt protocol).
- No LemonSqueezy tier strings — ever.

- Post-login redirect: /instrument only, never /daw.
- No hydrateFromToken() inside ProtectedRoute render.

12. Data Architecture & Schema

PostgreSQL Tables (Verified — [server/db/schema.ts](#))

Table	Key Columns	Status
users	id, email, hashedPassword, tier, isAdmin, createdAt	✓ Live
sessions	id, userId, startedAt, endedAt, metrics JSON	✓ Live
sessionMetrics	id, sessionId, aiActionsCount, timeSavedMs, acceptanceRate	✓ Live
subscriptions	id, userId, tier, stripeCustomerId, status	✓ Live
projects	id, userId, name, arrangementJSON, createdAt	✓ Live
samples	id, userId, filename, path, duration, sampleRate, bitDepth	✓ Live
presets	id, userId, name, type, data, isFactory	✓ Live
settings	id, userId, 17 columns, updatedAt	✓ Live
aiDecisionLog	id, sessionId, nodeId, actionType, trackId, inputConfidence, displayedConfidence, decision, outcome, latencyMs, timestamp	✓ Migration 0005 (2026-04-09)
effectPresetsTable	id, name, settings, createdAt, updatedAt	✓ Live
effectChainsTable	id, name, nodes, createdAt, updatedAt	✓ Live
djCuesTable	id, trackId, position, label, color, createdAt	✓ Live
waveformEditsTable	id, sampleId, editType, params, createdAt	✓ Live

AIDecisionLog — Key Schema (PRD §12 — Now Implemented)

```
id:                text PRIMARY KEY
sessionId:         text NOT NULL (FK to sessions)
nodeId:           text NOT NULL -- 'aiMixEngine' | 'transitionGraph'
actionType:       text NOT NULL -- 'gain_adjust' | 'eq_suggest' |
                  'transition_generate' | 'conflict_flag'
trackId:          text NULLABLE
inputConfidence:  real NOT NULL -- raw model confidence [0,1]
displayedConfidence: real NOT NULL -- after calibration
decision:         jsonb NOT NULL -- action-specific payload
outcome:          text NOT NULL -- 'auto_applied' | 'accepted' | 'rejected'
                  | 'ignored' | 'discarded'
latencyMs:        integer NOT NULL
```

timestamp: text NOT NULL

13. API Contract Specification

tRPC appRouter Shape (Verified — server/procedures.ts, 2026-04-09)

Router Key	File	Status
sessions	routers/sessions.ts	✓ Wired
sessionMetrics	routers/sessionMetrics.router.ts	✓ Wired
admin	routers/adminRouter.ts	✓ Wired
daw	routers/daw.ts	✓ Wired
subscription	routers/subscription.ts	✓ Wired
mixer	routers/mixer.router.ts	✓ Wired (added 2026-04-09)
dj	routers/dj.router.ts	✓ Wired (added 2026-04-09)
aiMix	routers/aiMix.router.ts	✓ Wired (added 2026-04-09)
projects	routers/index.ts (exported)	✓ Wired (added 2026-04-09)
presets	routers/index.ts (exported)	✓ Wired (added 2026-04-09)
settings	routers/index.ts (exported)	✓ Wired (added 2026-04-09)
ping	publicProcedure inline	✓ System health check

Key Procedure Contracts

ai.startPipeline

input: { sessionId: string; trackIds: string[] }
output: { pipelineId: string; nodeStatus: NodeStatus[] }

ai.submitSuggestionOutcome

input: { decisionId: string; outcome: 'accepted' | 'rejected' | 'ignored'; manualValueIfOverride?: number }
output: { logged: boolean; acceptanceRateUpdated: number }

sessions.endSession

input: { sessionId: string }
output: { metrics: SessionMetrics; exportPngUrl?: string }

14. Security & Compliance

Authentication Architecture

- Email/Password: bcrypt v6.0.0 → JWT (jsonwebtoken v9.0.3) → httpOnly cookie.
- Token Refresh: hydrateFromToken() race condition — FIXED.
- tRPC auth middleware: trpcAuth mount — FIXED and confirmed.

Authorization Layers

Layer	Method	Enforcement Point
Route-level	trpcAuth middleware	server/middleware/
Procedure-level	tRPC context validation	server/trpc.ts
Data-level	userId FK on every query	Drizzle ORM — all reads filtered
File access	Path traversal guard	Upload route — allowlist-validated
Subscription tier	Tier check in procedure	server/procedures.ts — all AI procedures

Rate Limiting

express-rate-limit v7.5.1 — 100 AI requests/minute/user
Message: 'AI request limit reached — upgrade to pro_artist for higher limits'

15. MVP Definition — Investor Checklist

Functional Requirements

Requirement	Status	Test
Load and play ≥8 simultaneous tracks	✔ Built	8-track test session
LLPTE node graph — animated, live inference	✔ Built	Node badge updates every 100ms
AI Auto-Leveling — active within 2 bars	✔ Built	Automated playback test
Ghost knobs visible on mixer strips	✔ Built	Visual QA
"Send to AI" context menu	✔ Built	Right-click on any clip
≥1 transition type (crossfade)	✔ Built	Execute and preview
Time Savings panel with real data	✔ Wired	SessionSummaryPanel in DAW.tsx
SessionChip in top nav	✔ Wired	Shows AI action count live
aiDecisionLog table	✔ Migration 0005	Schema + types + Zod schema
VU meters active on all 6 channel strips	✔ Built	Visual QA during playback
Effects rack: Reverb, Delay, Compressor	✔ Built	Knob interaction
Pro badge in top nav	✔ Built	Subscription tier check
TSC zero errors	✔ Verified	pnpm tsc --noEmit
Mix Suggestion System	☐ In Progress	MVP item 4 — frontend built, backend pending

Demo-Specific Requirements (PRD §26 — All Must Pass)

Requirement	Status	Failure Condition
inference 10ms badge updating live	✔	Static badge = demo failure
Transition Graph tooltip: edge count + tick time	✔	No tooltip = demo failure
Signal oscilloscope animated	✔	Static waveform = demo failure
Session summary exportable PNG	✔ Wired	Export fails = demo failure
aiDecisionLog writes live	☐	Acceptance rate shows 0 = weak demo

No audio dropouts in 20-minute demo		Any dropout = demo failure
No crashes		Any crash = demo failure

16. Risks & Mitigations

ID	Risk	Probability	Mitigation
R 1	AI feels gimmicky	Medium	Ship only features with measurable, audible impact. Auto-Leveling first.
R 2	LLPTE latency degrades under load	Medium	Confidence gating + auto-throttle. 8-track limit. Per-node tick monitoring live in UI.
R 3	User distrust of AI decisions	Medium	Ghost knobs + confidence scores + easy override. Never auto-apply below 0.65.
R 4	Demo crash in front of investor	Low	Pre-demo QA checklist — 17 items. Demo environment hardcoded pro_artist tier.
R 5	Audio artifacts from AI adjustments	Low	All gain changes via <code>AudioParam.setTargetAtTime()</code> — click-free by design.
R 6	Competitive response from Ableton/NI	Low	LLPTE is architectural — cannot be feature-flagged into legacy DAW.
R 7	Inference cost exceeds revenue	Medium	Local inference for MVP — zero cloud cost. Post-MVP: quantized models on-device.
R 8	Hygiene debt slows velocity	Medium	<code>r3_hygiene.py</code> patched: phantom dir exclusion, score formula uncapped, router key fixed (3 code fixes). 1 further fix flagged.

17. Post-MVP Roadmap

Phase 2 — Q3 2026 (6 months post-launch)

Theme: Depth + Community

Feature	Description	Rationale
Stem separation	AI-powered track separation from mixed audio	Unlocks mashup workflow for DJs
Custom AI preferences	User weights AI suggestions by genre/style	Reduces false positives for experienced users
Session history	Browse, replay, and fork past sessions	Retention — users revisit and improve past work
Hardware controller mapping	MIDI CC to mixer/effects parameters	Bridges R3 to existing DJ hardware
Collaboration (beta)	Real-time multi-user session sharing	Unlocks pro_artist tier use cases

Phase 3 — Q1 2027 (12 months post-launch)

Theme: Platform + B2B

Feature	Description	Rationale
LLPTE API (public)	tRPC endpoints exposed for third-party integrations	B2B revenue — DAW plugins, hardware OEM
Mobile companion	iOS/Android remote transport control	Extend studio to mobile without rebuilding core
Advanced analytics	Per-genre AI calibration, personalized baseline	Model improves with user data flywheel
Label/studio tier	Multi-user workspace with admin controls	Enterprise contract opportunity

Phase 4 — Q3 2027 (24 months post-launch)

Theme: AI Model + Market Expansion

Feature	Description	Rationale
R3 model fine-tuning	User fine-tunes LLPTE on own mix history	True personalization — highest retention driver
Generative transitions	AI generates entirely new transition content	Category-defining feature
Hardware partnerships	R3 software bundled with DJ controllers	Distribution at hardware scale

Marketplace	Community-shared AI presets and transition styles	Network effects — platform moat
-------------	---	---------------------------------

18. Future Development & Upgrade Pathways

This section documents technical upgrade pathways not covered in the Phase roadmap. These are engineering-level decisions with long-horizon implications for the LLPTE architecture, infrastructure, and product surface.

18.1 LLPTE Model Upgrades

Current State

LLPTE uses heuristic confidence scoring — rules-based thresholds calibrated against 50-session alpha baselines. No trained model is deployed. aiDecisionLog (migration 0005) now captures labeled outcomes, beginning the data flywheel.

Upgrade Path

- Phase 1 (now → 100K decisions): Accumulate labeled data via aiDecisionLog.outcome field. Every accepted/rejected suggestion is a training signal.
- Phase 2 (100K decisions): Train first lightweight classification model on gain_adjust decisions. Deploy as WASM module in llpte-ai package — zero latency regression.
- Phase 3 (1M decisions): Per-user calibration. Rolling acceptance rate from mv_ai_acceptance_rates materialized view feeds confidence threshold adjustment per user.
- Phase 4 (post-Series A): Full per-genre model fine-tuning. LLPTE API exposed publicly — third parties can submit labeled audio decisions to improve the shared model.

18.2 Infrastructure Upgrades

Current

Railway (backend) + Vercel (frontend). Local inference only — zero cloud compute cost. PostgreSQL managed. S3/R2 for audio file storage (server/lib/storage-s3.ts via Proxy pattern).

Upgrade Path

- 10K users: Add mv_user_session_averages and mv_ai_acceptance_rates materialized views to PostgreSQL — currently missing, required for accurate Time Savings baseline calculation.
- 25K users: Migrate inference to dedicated GPU worker pool. AudioWorklet WASM modules remain local — only model inference moves to cloud. Estimated cost: \$0.0008 per inference at 25K users = ~\$600/month.
- 100K users: Read replicas for PostgreSQL. CDN for waveform peak data. OffscreenCanvas compute shaders for WebGPU waveform rendering already implemented — no change needed.
- Post-\$5M ARR: On-premise Docker deployment option (Dockerfile + docker-compose.yml already in repo). Required for enterprise/label tier SLA guarantees.

18.3 Audio Engine Upgrades

Current

Web Audio API + AudioWorklet (WASM module) + SharedArrayBuffer ring buffer. 8-track ceiling enforced at LLPTE layer.

Upgrade Path

- Track count expansion: LLPTE active edge limit is 2000 (current: 847). 16-track support achievable without architecture change. 32-track requires transitionGraph optimization pass.
- SIMD acceleration: Wasm SIMD already planned in architecture. Float32 SIMD for DSP operations. Target: 3ms processing per frame (currently undocumented baseline).
- WebGPU compute shaders: OffscreenCanvas already implemented for waveform rendering. Extend to FFT computation — moves spectralAnalyzer workload to GPU. Eliminates primary CPU bottleneck at high track counts.
- Stem separation (Phase 2): Demucs v4 quantized WASM model. Runs in AudioWorklet. 4-stem output (vocals, drums, bass, other). Estimated inference: 8–12 seconds per track on consumer GPU.

18.4 Frontend Architecture Upgrades

Current

React 18 + Vite 5 + Wouter + Zustand. DAW.tsx is ~1,800 lines — single-file monolith. Multiple phantom directories (client/client, client/hooks, client/stores) exist outside client/src/ and create import graph ambiguity.

Upgrade Path

- Immediate: Resolve phantom directory imports. client/src/store/auth-store.ts imports from ../hooks/authStore (phantom). Consolidate all stores under client/src/stores/, all hooks under client/src/hooks/.
- Near-term: Split DAW.tsx into feature modules. Suggested structure: DAW/TransportBar, DAW/ArrangementView, DAW/MixerSection, DAW/LLPTEGraph, DAW/EffectsRack. Each under client/src/features/daw/.
- Medium-term: React Server Components for sample browser and file tree. Reduces initial bundle. Vite 5 already supports RSC — no build tool change required.
- Long-term: Move to React 19 when stable. Concurrent features (useTransition, useDeferredValue) benefit the heavy FFT visualization renders without audio thread blocking.

18.5 Monetization Upgrades

Current

Stripe v20.4.1. Three tiers: explorer (free) / creator (\$20/month) / pro_artist (\$60/month). Subscription router wired. aiDecisionLog provides acceptance rate data required for upgrade prompt triggers.

Upgrade Path

- Contextual upgrade prompts: Wire aiDecisionLog acceptance rate to upgrade prompt logic. Trigger "upgrade to creator" when free user hits 5 AI uses in a session.
- Annual billing: Add annual_creator (\$192/year, 20% discount) and annual_pro_artist (\$576/year) Stripe price IDs. Server-side only — no schema change required.
- Usage-based billing (post-\$1M ARR): Charge per AI inference above tier limits. Stripe Meters API (introduced 2024). Requires inference counting middleware in sessionMetrics.service.ts.
- B2B API licensing (Phase 3): LLPTE tRPC endpoints exposed under /api/v1/llpte. Per-call pricing via Stripe Meters. Target customers: DAW plugin developers, hardware OEMs, content platforms.

18.6 Codebase Hygiene Upgrade Path

Current State (2026-04-12 — v4.1 update)

TSC errors	0
Routers wired	11/11
Live any violations	5 remaining: routes/presets.ts × 4 (Drizzle insert casts), audio-analysis.ts × 1 (eslint-disable exemption — no .d.ts for node-web-audio-api)
console.log violations	5 remaining in server/index.ts:300-308 (separate hard guard — replace with structured logger)
Phantom directories	9 remaining (client/client, client/hooks, client/components, client/stores, client/src/hook, client/src/context, client/src/contexts, client/src/store [LIVE — has imports], db/schema)
Hygiene script bugs fixed	3 code-fixed (phantom dir exclusion, score formula uncapped, router key subscription), 1 flagged (PRD items hardcoded — future fix)
Migration applied (Railway)	PENDING — 0005_overjoyed_gambit.sql

Path to 90/100 Hygiene Score

- P0: Apply migration 0005 to Railway production DB.
- ✓ COMPLETED 2026-04-12: MultitrackView (/mixer) — double transport bar eliminated via hideTransport={true} prop. Acid-techno theme (#0a0a0a / #a3e635 / IBM Plex Mono) applied to mixer page. TSC: 0 errors.
- ✓ COMPLETED 2026-04-12: SKILLS.md created — 22 engineering patterns and failure modes documented in docs/SKILLS.md.
- ✓ COMPLETED 2026-04-12: DEMO_CHECKLIST.md created — full pre-demo QA checklist in docs/DEMO_CHECKLIST.md.
- P0: Wire aiDecisionLog writes into session-metrics.service.ts — enables live acceptance rate tracking.
- P1: Fix routes/presets.ts — 4 Drizzle as any casts — replace with typed InsertEffectPreset / InsertEffectChain.

- P2: Replace console.log in server/index.ts:300-308 with structured logger (morgan already installed).
- P3: Consolidate phantom directories — merge client/client, client/hooks, client/components, client/stores into client/src/ canonical locations.
- P4: Create mv_user_session_averages and mv_ai_acceptance_rates materialized views in Drizzle migration 0006.
- P5: Fix vitest root config — add include pattern for packages/*/tests/*.test.ts to get accurate test count for PRD §0/§22/§25.

19. Go-to-Market Strategy

Channel 1: Content-Led (Primary — 0 to \$500K ARR)

- "Before AI / After AI" side-by-side audio comparisons (60 seconds)
- LLPTE node graph explanations — "This is what R3 actually does to your tracks"
- Time Savings reveal videos — "R3 just told me I mixed 42% faster"
- Distribution: Twitter/X, TikTok, YouTube Shorts, r/DJs, r/WeAreTheMusicMakers
- KPI: 100 signups per viral content piece at >100K views

Channel 2: Community (Primary — 0 to \$500K ARR)

- DJTT (DJ Tech Tools) — editorial relationship + sponsored content
- Discord DJ communities — authentic participation
- Creator economy newsletters (Matt Gielen, Paddy Galloway)
- KPI: 500 alpha users from community channels in first 90 days

Channel 3: Hardware Partnerships (Medium-term — \$500K to \$3M ARR)

Bundle R3 creator tier (3-month trial) with mid-tier DJ controller purchases via OEM agreement. Pioneer DJ, Rane, Denon DJ collectively ship 3M+ units annually. Target: 10% of 500K relevant units = 50K new users. At 10% conversion: 5,000 subscribers = \$100K MRR.

20. Financial Model & Valuation Gates

Cost Structure (Pre-Scale)

Infrastructure (Railway)	\$200/month
PostgreSQL (managed)	\$100/month
AI inference (local — MVP)	\$0 — no cloud inference cost
Storage (S3/R2)	\$50/month
Stripe fees	2.9% + \$0.30 per transaction
Error monitoring	\$26/month
Analytics	\$50/month
Total fixed/semi-fixed	~\$450/month

Valuation Gates

State	Valuation Range	Gap from Current
Current — technical asset only	\$180K–\$400K	Baseline
Working demo + 50 beta users	\$800K–\$2.5M pre-seed	4–6 weeks execution
≥65% AI acceptance rate confirmed	\$3–6M seed	Requires aiDecisionLog writes live
\$120K ARR (500 paying users)	\$4.8–9.6M (40–80× ARR, typical pre-seed SaaS)	12 months post-launch
\$1M ARR	\$15–30M Series A (15–30× ARR)	24 months post-launch

The current build state — LLPTE implemented, UI shipped, 11 routers wired, auth solid, database schema complete including aiDecisionLog — represents approximately \$300–600K of technical value as a pure asset. The gap between current state and seed metrics is approximately 4–6 weeks of focused execution.

21. Investor Demo Script (Updated — 12 Minutes)

Setup (30 minutes before)

- Run pre-demo QA checklist — 17 items confirmed (see §15)
- Load demo session: 8 pre-selected tracks, 128 BPM, key of A minor
- Confirm Pro badge visible in top nav
- Confirm LLPTE node graph visible and animated (all 5 connector lines moving)
- Confirm SessionChip in top nav is inactive (session not yet started)
- Dim room lights — screen is the only light source

Demo Flow

Minute	Action	Key Visual
0–1	Press Space. Say: "Watch the bottom-right corner."	LLPTE node graph animates. SessionChip shows ● Live.
1–2	Within 2 bars: toast fires "AI Mix updated — confidence 94%"	Ghost knobs appear on all channel strips simultaneously.
2–4	Click Accept All. Demonstrate one manual override.	Knobs animate to new positions. Override logged.
4–6	Suggestion panel fires: "Cut low-end on SYNTH LEAD below 80Hz"	"The AI caught a frequency clash in 6 seconds."
6–8	Right-click VOCAL CHOP → Send to AI → Key-Match Crossfade → Preview → Accept	Violet overlay. Camelot scoring. Transition committed.
8–10	Hover Transition Graph node — tooltip: "847 active edges last tick 0.8ms"	"10ms inference. Industry best is 15ms." Pause.
10–12	Press Space to stop. SessionSummaryPanel appears.	"42% faster. 18 minutes saved. 34 AI actions. 71% acceptance." Export PNG. Pause. "New category."

22. Glossary

LLPTE	Low-Latency Processing Transition Engine — R3's named real-time AI pipeline, 6 TypeScript monorepo packages
inputRouter	LLPTE node 1 — routes multi-track audio into pipeline, normalizes format
spectralAnalyzer	LLPTE node 2 — FFT, RMS, LUFS, true peak measurement per frame
aiMixEngine	LLPTE node 3 — inference engine, produces gain/EQ decisions with confidence scores
transitionGraph	LLPTE node 4 — precomputes and executes transition curves; 847 active edges in demo
outputBus	LLPTE node 5 — final signal delivery and monitoring
Ghost knob	UI element showing where AI wants to move a mixer parameter — translucent violet over actual knob
Confidence score	AI's self-assessed probability a decision is correct — gate at 0.65 for auto-apply
Confidence gating	Logic determining auto-apply (≥ 0.65), suggestion (≥ 0.40), or discard (< 0.40)
aiDecisionLog	PostgreSQL table logging every AI decision with outcome — the data flywheel. Migration 0005.
SessionChip	Top nav component showing live session state and AI action count. Wired 2026-04-09.
SessionSummaryPanel	Full-screen overlay on session end showing time savings, acceptance rate, export. Wired 2026-04-09.
Send to AI	Right-click context menu action routing audio clip into LLPTE pipeline for transition generation
Camelot wheel	Harmonic compatibility scoring for DJ mixing — used for key-match crossfade confidence
explorer	Free tier (correct DB value — NOT "free" or "Starter")
creator	Primary paid tier at \$20/month (NOT "Pro")
pro_artist	Studio tier at \$60/month (NOT "Studio")
Wire.txt protocol	Engineering discipline: Files Read → Findings → Changes → Remaining Ambiguities. No write without read first.
LUFS	Loudness Units Full Scale — broadcast-standard loudness measurement (ITU-R BS.1770-4)
Zinc-950	Tailwind CSS color token — darkest background in R3 design system (#09090b)
r3_hygiene.py	Python-based 10-phase monorepo audit script. 3 code bugs fixed 2026-04-09 (phantom dir exclusion, score formula, router key). 1

	further fix flagged.
--	----------------------

FINAL PRINCIPLE

This document is a contract, not a wish list.

Everything grounded in what is built, what is partially built, or what has a clear engineering path to completion.

Prove that AI reduces the time and skill required to create a clean, professional mix — in real time.

The LLPTE pipeline makes that proof visible. The Time Savings panel makes it quantifiable. The ghost knobs make it trustworthy. The 10ms inference badge makes it undeniable.

R3 v4 — PRD v4.1 — 2026-04-12